
SCS 1205

Computer Systems

Kasun Gunawardana

E-mail: kgg



University of Colombo School of Computing

Memory Organization

Primary & Secondary Memory

- Primary Memory
 - Cache Memory
 - Main Memory
- Secondary Memory
 - Magnetic Disks
 - Magnetic Tapes
 - USB Sticks

Types of Memory

- Read Only Memory (ROM)
 - Non-volatile
 - Variants
 - PROM (Programmable Read Only Memory)
 - EEPROM (Electrically Erasable Programmable Read Only Memory)
- Random Access Memory (RAM)
 - Volatile
 - Variants
 - Dynamic RAM (DRAM)
 - Static RAM (SRAM)

Flash Memory

- Floating-gate metal–oxide–semiconductor field-effect transistors
- Non-volatile
- Read and Write
- Common Usages
 - Memory Cards
 - USB Flash Drives
 - Solid State Drives

Solid State Devices - SSD

- Typically based on Flash Memory
- Compared to HDD
 - More resistant to physical shock
 - Run silently
 - Faster access time
- Not appropriate for long term storage
 - Without power it depletes charge over time

3D XPoint

- Newly developed memory technology
- Read more about it...
 - Start with https://en.wikipedia.org/wiki/3D_XPoint


Memory Hierachy

Organization of computer storage components based on different attributes.

- Access time/ Latency
- Cost
- Size
- Distance from CPU

Memory Hierarchy

Computer Memory Hierarchy



Typical Data Access Pattern

- CPU requests data from its nearest storage
 - Cache
- If not found, then request goes to the Main Memory
- If not found, then request goes to the HDD
- Once the data is found it will be loaded into the Main Memory and then into Cache

The CPU – Memory Gap

- CPU Cycle time decreases rapidly
- Memory access times cannot match the CPU speed
 - HDD > SSD > DRAM > SRAM >> CPU
- CPU - Memory gap affects the system performance
- **“Principle of Locality”** is used to bridge the gap

Principle of Locality


Tendency of CPU to access the same set of memory locations repeatedly over a short period of time.

Programs tend to refer data and instructions with addresses near or equal to those they have referred recently.

Principle of Locality


- Spatial Locality:

Items with nearby addresses tend to be referenced close together in time



- Temporal Locality:

Recently referenced items are likely to be referenced again in the near future



Reading...

- https://en.wikipedia.org/wiki/Locality_of_reference
- <https://stackoverflow.com/questions/7638932/what-is-locality-of-reference>
- <https://www.geeksforgeeks.org/locality-of-reference-and-cache-operation-in-cache-memory/>

Cache Memory

- A smaller and faster memory space that keeps data of a larger and slower memory space.
- Reside in CPU chip
- Typical characteristics
 - Transient – Highly volatile
 - May keep frequently referred data
 - May keep data that would be referred in future

Caching...



[This Photo](#) by Unknown Author is licensed under [CC BY-NC](#)

Cache Hierarchy

- Cache memories are typically SRAMs
- Per unit cost is high
- Faster and larger cache space will be expensive
- As a solution hierarchy of caches is used
 - Different Levels of Caches (i.e. L1, L2, L3)

Reading..

- <https://stackoverflow.com/questions/3699582/what-is-the-difference-between-l1-cache-and-l2-cache>

Cache Performance

- **Cache Hit:** Requested data for processing by a CPU is found in the cache memory.
 - **Cache Miss:** Requested data is not found.
-
- **Cache Hit Ratio:** The probability of finding requested data in the cache.
 - **Cache Miss Ratio** = $1 - \text{Cache Hit Ratio}$

Cache Performance (Cont.)

- Effective Access Time (EAT)
 - Average time it takes for CPU to access data
- $EAT = HR \times T_C + (1 - HR) \times T_M$

HR – Hit Ratio

T_C – Cache Access Time in a cache hit

T_M – Data access time in a cache miss

Key Design Decisions of a Cache

- 1. Cache Volume (Capacity, Size)
- 2. Write Policy
 - Write Back
 - Write Through
- 3. Mapping Scheme
 - Direct Mapped
 - Set Associative
 - Line replacement policies

2. Cache – Write Policy

Main Memory	
0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
...	...



Cache Memory	
2	c
3	d
10	k
11	l

What if CPU
changes one of the
values in cache ?

2. Cache – Write Policy (Cont.)

Main Memory	
0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
...	...



Cache Memory	
2	c
3	d
10	k
11	l

How shall we ensure consistency and reliability of memory content?

2. Cache – Write Policy (Cont.)

- Cache is a redundant storage.
- Same data appears in two locations.
- Data maybe shared across processes.
- Maintaining consistency is a challenge when the same data appears in multiple locations.
- Consistency in Cache and Main Memory.
- Inconsistency - The data is altered (write).
- Cache needs a write policy to ensure consistency.

Write Policy – Write Through

Whenever write operation happens, it is performed on both cache and main memory synchronously.

- Comparatively simple policy to be implemented.
- Bus traffic will be increased.

Write Policy – Write Back

Write operations on a data is performed only on cache. But whenever the data is evicted from the cache it is written to the main memory.

- Lazy writes.
- Complex to implement.

Reading Exercise...

Cache coherence

- https://en.wikipedia.org/wiki/Cache_coherence

3. Mapping Schemes

Main Memory

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
...	...



Cache Memory

2	c
3	d
10	k
11	l

How shall we map memory identifiers to the cache ?

3. Mapping Schemes (Cont.)


- The CPU refers data or instruction by its addresses.
- The data is first searched in Cache.
- Thus, cached data should be able to access using addresses.
- But cache is smaller compared to the main memory.
- Cache memory locations cannot be given main memory addresses.
- A mapping scheme is required.

Mapping Schemes (Cont.)

- Direct Mapped
- Fully Associative
- Set Associative

Overview of a Cache


Main Memory	
0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
...	...



Cache Memory	

Overview of a Cache

Main Memory	
0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
...	...



Spatial Locality
by bringing a block
of data

Cache Line (Cache Block)

- Data is transferred between memory and cache in blocks of fixed size, called **Cache Lines** (cache blocks).
- When a **Cache Line** is copied from the memory into the cache, a **Cache Entry** is created.
- The **Cache Entry** will include the copied data as well as the requested memory location identifier (called a tag).
- A single cache line may contain multiple bytes/words of data.
- **Amount of data** in a **Cache Entry** and a **Cache Line** are always equivalent in size.

Next...

Cache Mapping Schemes

- Direct Mapped Cache
-

Computer Systems

Kasun Gunawardana

E-mail: kgg




University of Colombo School of Computing

Direct Mapped Cache

Direct Mapped Cache

- A memory is identified as blocks of fixed size.
- Each block has a pre-defined cache entry position.
- A block in the memory will always map to the same cache entry position.
- The number of blocks in the main memory is larger than the maximum number of cache entries.
- Thus, several memory blocks are mapped to the same cache entry position.
- Tag is used to identify the address of a mapped data.


Direct Mapped Cache



- Byte addressable main memory
- Total of 1024 Bytes in the memory
- 2 Bytes per cache line
- Cache volume = 4 Bytes


0	
1	
2	
3	

Direct Mapped Cache



- Byte addressable main memory
- Total of 1024 Bytes in the memory
- 2 Bytes per cache line
- Cache volume = 4 Bytes

Direct Mapped Cache





- Byte addressable main memory
- Total of 1024 Bytes in the memory
- **2 Bytes per cache line**
- Cache volume = 4 Bytes

0	
1	
2	
3	

Direct Mapped Cache


- Byte addressable main memory
- Total of 1024 Bytes in the memory
- 2 Bytes per cache line
- Cache volume = 8 Bytes



Direct Mapped Cache Index




- Byte addressable main memory
- Total of 1024 Bytes in the memory
- 2 Bytes per cache line
- Cache volume = 8 Bytes



Direct Mapped Cache Offset




- Byte addressable main memory
- Total of 1024 Bytes in the memory
- 2 Bytes per cache line
- Cache volume = 8 Bytes



Direct Mapped Cache Tag




- Byte addressable main memory
- Total of 1024 Bytes in the memory
- 2 Bytes per cache line
- Cache volume = 8 Bytes



Direct Mapped Cache Tag

a	0000000 00 0
b	0000000 00 1
c	0000000 01 0
d	0000000 01 1
e	0000000 10 0
f	0000000 10 1
g	0000000 11 0
h	0000000 11 1
i	0000001 00 0
j	0000001 00 1
...
yyy	1111111 11 0
zzz	1111111 11 1


- Byte addressable main memory
- Total of 1024 Bytes in the memory
- 2 Bytes per cache line
- Cache volume = 8 Bytes



Direct Mapped Cache Tag

a	0000000 00 0
b	0000000 00 1
c	0000000 01 0
d	0000000 01 1
e	0000000 10 0
f	0000000 10 1
g	0000000 11 0
h	0000000 11 1
i	0000001 00 0
j	0000001 00 1
...
yyy	1111111 11 0
zzz	1111111 11 1


- Minimum number of bits required for
 - Offset
 - Index
 - Tag



Direct Mapped Cache – Example

a	0000000	00	0
b	0000000	00	1
c	0000000	01	0
d	0000000	01	1
e	0000000	10	0
f	0000000	10	1
g	0000000	11	0
h	0000000	11	1
i	0000001	00	0
j	0000001	00	1
...		
yyy	1111111	11	0
zzz	1111111	11	1


- CPU wants to read address 1
 - 1 = 0000000 00 1



Direct Mapped Cache – Example

a	0000000 00 0
b	0000000 00 1
c	0000000 01 0
d	0000000 01 1
e	0000000 10 0
f	0000000 10 1
g	0000000 11 0
h	0000000 11 1
i	0000001 00 0
j	0000001 00 1
...
yyy	1111111 11 0
zzz	1111111 11 1


- CPU wants to read address 1
- 1 = 0000000 **00 1**



Direct Mapped Cache – Example

a	0000000 00 0
b	0000000 00 1
c	0000000 01 0
d	0000000 01 1
e	0000000 10 0
f	0000000 10 1
g	0000000 11 0
h	0000000 11 1
i	0000001 00 0
j	0000001 00 1
...
yyy	1111111 11 0
zzz	1111111 11 1


- CPU wants to read address 1
- 1 = 0000000 **00 1**



Direct Mapped Cache – Example

a	0000000 00 0
b	0000000 00 1
c	0000000 01 0
d	0000000 01 1
e	0000000 10 0
f	0000000 10 1
g	0000000 11 0
h	0000000 11 1
i	0000001 00 0
j	0000001 00 1
...
yyy	1111111 11 0
zzz	1111111 11 1


- CPU wants to read address 4
- $4 = 0000000 \boxed{10} \ 0$



Direct Mapped Cache – Example

a	0000000 00 0
b	0000000 00 1
c	0000000 01 0
d	0000000 01 1
e	0000000 10 0
f	0000000 10 1
g	0000000 11 0
h	0000000 11 1
i	0000001 00 0
j	0000001 00 1
...
yyy	1111111 11 0
zzz	1111111 11 1


- CPU wants to read address 8
- $8 = 0000001 \boxed{00} \ 0$



Direct Mapped Cache – Example

a	0000000 00 0
b	0000000 00 1
c	0000000 01 0
d	0000000 01 1
e	0000000 10 0
f	0000000 10 1
g	0000000 11 0
h	0000000 11 1
i	0000001 00 0
j	0000001 00 1
...
yyy	1111111 11 0
zzz	1111111 11 1


- CPU wants to read address 8
- $8 = 0000001 \boxed{00} \ 0$



Direct Mapped Cache – Example

a	0000000 00 0
b	0000000 00 1
c	0000000 01 0
d	0000000 01 1
e	0000000 10 0
f	0000000 10 1
g	0000000 11 0
h	0000000 11 1
i	0000001 00 0
j	0000001 00 1
...
yyy	1111111 11 0
zzz	1111111 11 1

- CPU wants to read address 8
- $8 = 0000001 \boxed{00} \ 0$



Direct Mapped Cache - Searching

- Use the Low Order Address Bits (LOAB) of an address to index into the Cache Tag Memory.
- Use comparator to find whether the corresponding Tag matches with Higher Order Address Bits of the address.
- If the Tag matches then check for the valid bit.
 - Valid Bit =1 then Cache Hit
 - Valid Bit =0 then Cache Miss
- If the Tag not found then it is Cache Miss

Exercise

- 2 MB Direct-Mapped cache
- Cache line size is 32 KB
- Byte addressable 1GB memory



1. What are lengths of tag, index and offset?

Exercise

- 2 MB Direct-Mapped cache
- Cache line size is 32 KB
- Word addressable 1GB memory
- Word size is 32 bits



1. What are lengths of tag, index and offset?

Exercise

- 2 MB Direct-Mapped cache
- Cache line size is 32 KB
- Word addressable 1GB memory
- Word size is 32 bits

- Cache Line = 32KB
- Word = Smallest Unit which has an address
- Offset = How many addressable units within a cache line?

$$\frac{32KB}{32 \text{ bits}} = \frac{32 \times 2^{10} \text{ Bytes}}{4 \text{ Bytes}} = \frac{2^{15} \text{ Bytes}}{2^2 \text{ Bytes}}$$



$$2^{13}$$

- How many addressable units within a cache line? 2^{13}
- Therefore, we need at least 13 bits for Offset

Exercise

- 2 MB Direct-Mapped cache
- Cache line size is 32 KB
- Word addressable 1GB memory
- Word size is 32 bits

- Cache Line = 32KB
- Index = How many cache entries can we create ?

$$\frac{2MB}{32\text{ KB}} = \frac{2 \times 2^{20} \text{ Bytes}}{2^5 \times 2^{10} \text{ Bytes}} = \frac{2^{21} \text{ Bytes}}{2^{15} \text{ Bytes}}$$

$$2^6$$



- How many cache entries can we create? 2^6
- Therefore, we need at least 6 bits for Index

Exercise

- 2 MB Direct-Mapped cache
- Cache line size is 32 KB
- Word addressable 1GB memory
- Word size is 32 bits



- Minimum address length ??
- Main memory = 1GB
- Word size = 32 bits = 4 Bytes
- How many addresses do we have ?

$$\frac{1GB}{32 \text{ bits}} = \frac{2^{30} \text{ Bytes}}{4 \text{ Bytes}} = 2^{28}$$

- Therefore, we need at least 28 bits for address



Homework..

- Find advantages and disadvantages of Direct Mapped Cache

Next...

Cache Mapping Schemes

- Set Associative Cache
-

Computer Systems

Kasun Gunawardana

E-mail: kgg



University of Colombo School of Computing

Set Associative Cache

Fully Associative Cache

- Direct Mapped Cache affirms strict mapping rule.
- Fully Associative Cache is the opposite.
 - No restrictions.
 - Cache lines can be placed on anywhere available.
- More flexible cache mapping scheme.
- Searching is expensive.




Set Associative Cache

- Hybrid model of Fully Associative and Direct Mapped caches.
- Cache is divided into equal sized sets.
- Cache line can be mapped into only one set (Directly mapped to a set).
- Within the mapped set, cache line can be placed anywhere available (Fully Associative within a set).

If a set in the cache can accommodate N number of cache lines, then it is called “N-Way Set Associative Cache”.

N-Way Set Associative


- If is a cache line and segmentation shows the sets, recognize N in following cache arrangements.



N-Way Set Associative


- If [] is a cache line and segmentation shows the sets, recognize N in following cache arrangements.

1-way Associative



8 sets

??




4 sets

??



2 sets

??




1 set


N-Way Set Associative

- If [] is a cache line and segmentation shows the sets, recognize N in following cache arrangements.


1-way Associative




2-way Associative



4-way Associative



8-way Associative



8 sets

4 sets


2 sets

1 set


N-Way Set Associative

- If [] is a cache line and segmentation shows the sets, recognize N in following cache arrangements.

1-way Associative




2-way Associative



4-way Associative



8-way Associative



8 sets

4 sets

2 sets

1 set

Example: 2-Way Set Associative

	1 Byte
A	00000000 0 0
B	00000000 0 1
C	00000000 1 0
D	00000000 1 1
E	00000001 0 0
F	00000001 0 1
G	00000001 1 0
H	00000001 1 1
I	00000010 0 0
J	00000010 0 1
...	
...	
....	

- Byte addressable 2-way set associative cache
- Total of 1024 Bytes in the memory
- 2 Bytes per cache line
- Cache volume = 8 Bytes
- How do we arrange the cache layout?

Example:

- Byte addressable 2-way set assoc.
- Total of 1024 Bytes in the memory
- 2 Bytes per cache line
- Cache volume = 8 Bytes

A	00000000 0 0
B	00000000 0 1
C	00000000 1 0
D	00000000 1 1
E	00000001 0 0
F	00000001 0 1
G	00000001 1 0
H	00000001 1 1
I	00000010 0 0
J	00000010 0 1
...	
...	
....	

- 2-Way set associative – Within a single set there will be room for two cache lines.
- Thus the set size = 4 Bytes
- Since the cache volume is 8 Bytes, there will be only two sets.
- A cache line placed in a set is identified by a Tag
- Within a cache entry, each unit of memory is identified by the offset

Example:

- Byte addressable 2-way set assoc.
- Total of 1024 Bytes in the memory
- 2 Bytes per cache line
- Cache volume = 8 Bytes

A	00000000 0 0
B	00000000 0 1
C	00000000 1 0
D	00000000 1 1
E	00000001 0 0
F	00000001 0 1
G	00000001 1 0
H	00000001 1 1
I	00000010 0 0
J	00000010 0 1
...	
...	
....	

- A cache line will be mapped to one of two sets.
- We need set identifier
- Within mapped set, there will be cache entries.
- Tag is required to identify the address.
- Offset will be used to refer the exact memory unit within the cache entry.

Example:

- Byte addressable 2-way set assoc.
- Total of 1024 Bytes in the memory
- 2 Bytes per cache line
- Cache volume = 8 Bytes

A	00000000 0 0
B	00000000 0 1
C	00000000 1 0
D	00000000 1 1
E	00000001 0 0
F	00000001 0 1
G	00000001 1 0
H	00000001 1 1
I	00000010 0 0
J	00000010 0 1
...	
...	
....	




Example:

- Byte addressable 2-way set assoc.
- Total of 1024 Bytes in the memory
- 2 Bytes per cache line
- Cache volume = 8 Bytes

A	00000000 0 0
B	00000000 0 1
C	00000000 1 0
D	00000000 1 1
E	00000001 0 0
F	00000001 0 1
G	00000001 1 0
H	00000001 1 1
I	00000010 0 0
J	00000010 0 1
...	
...	
....	

You may remember, DMC made cache replacement when the CPU referred addresses 1 and 8 in order. What will happen here if the CPU refers the same?




Example:

- Byte addressable 2-way set assoc.
- Total of 1024 Bytes in the memory
- 2 Bytes per cache line
- Cache volume = 8 Bytes

A	00000000 0 0
B	00000000 0 1
C	00000000 1 0
D	00000000 1 1
E	00000001 0 0
F	00000001 0 1
G	00000001 1 0
H	00000001 1 1
I	00000010 0 0
J	00000010 0 1
...	
...	
....	

CPU refers address 1




Example:

- Byte addressable 2-way set assoc.
- Total of 1024 Bytes in the memory
- 2 Bytes per cache line
- Cache volume = 8 Bytes

A	00000000 0 0
B	00000000 0 1
C	00000000 1 0
D	00000000 1 1
E	00000001 0 0
F	00000001 0 1
G	00000001 1 0
H	00000001 1 1
I	00000010 0 0
J	00000010 0 1
...	
...	
....	

CPU refers address 1

- Mapping set is 0
- Can place anywhere within the set 0
(Assuming the cache is empty at the beginning)




Example:

- Byte addressable 2-way set assoc.
- Total of 1024 Bytes in the memory
- 2 Bytes per cache line
- Cache volume = 8 Bytes

A	00000000 0 0
B	00000000 0 1
C	00000000 1 0
D	00000000 1 1
E	00000001 0 0
F	00000001 0 1
G	00000001 1 0
H	00000001 1 1
I	00000010 0 0
J	00000010 0 1
...	
...	
....	

CPU refers address 8

- Mapping set is 0
- Still it can be placed. Because set 0 can accommodate another cacheline.




Example:

- Byte addressable 2-way set assoc.
- Total of 1024 Bytes in the memory
- 2 Bytes per cache line
- Cache volume = 8 Bytes

A	00000000 0 0
B	00000000 0 1
C	00000000 1 0
D	00000000 1 1
E	00000001 0 0
F	00000001 0 1
G	00000001 1 0
H	00000001 1 1
I	00000010 0 0
J	00000010 0 1
...	
...	
....	

CPU refers address 8

- Mapping set is 0
- Still it can be placed. Because set 0 can accommodate another cacheline.




Example:

- Byte addressable 2-way set assoc.
- Total of 1024 Bytes in the memory
- 2 Bytes per cache line
- Cache volume = 8 Bytes

A	00000000 0 0
B	00000000 0 1
C	00000000 1 0
D	00000000 1 1
E	00000001 0 0
F	00000001 0 1
G	00000001 1 0
H	00000001 1 1
I	00000010 0 0
J	00000010 0 1
...	
...	
....	

But, if CPU refers address 4 or 5 there will be an eviction.

- Mapping set is 0
- One cache entry must be replaced with new cacheline



In this circumstance, which cache entry should be evicted ?


Example:

- Byte addressable 2-way set assoc.
- Total of 1024 Bytes in the memory
- 2 Bytes per cache line
- Cache volume = 8 Bytes

A	00000000 0 0
B	00000000 0 1
C	00000000 1 0
D	00000000 1 1
E	00000001 0 0
F	00000001 0 1
G	00000001 1 0
H	00000001 1 1
I	00000010 0 0
J	00000010 0 1
...	
...	
....	

But, if CPU refers address 4 or 5 there will be an eviction.

- Mapping set is 0
- One cache entry must be replaced with new cacheline



In this circumstance, which cache entry should be evicted ?

Cache Replacement Policies

- When implementing set associative caches, we need to affirm a cache replacement policy.
 - Which cache entry should be evicted when the set is fully occupied?
- This decision should be taken aligned with the Locality of reference that we need to exploit.

Common Policies

- Least Recently Used
 - The cache entry which is unused for the longest time will be evicted.
- First In First Out
 - The cache entry which has the longest period of stay will be evicted.
- Least Frequently Used
 - The cache line with least access frequency will be evicted.
- Random

Homework...

- Discuss advantages and disadvantages of those cache replacement policies.

Exam

- Try to answer past paper questions.
- Discuss doubts and clarify with your colleagues.
- Exam Paper – 2 Hour paper –
 - 4 compulsory questions
 - A question from each section of the course

Thank You..!
