

# SCS 1214: Operating Systems

Dr. Dinuni Fernando, PhD

Based on “Virtual Machines” ,Smith and Nair,  
Chapter 1  
Chapter 7, Andrew Tanenbaum's Operating systems  
concepts


**Virtual Machines**



# Virtualization

- Makes a real system appear to be a set of virtual systems.
- Types of Virtualizations
  1. One-to-many virtualization
    - E.g. one physical machine may appear as multiple virtual machines
    - one physical disk may look like multiple virtual disk
    - one physical network may look like multiple virtual networks
  2. Many-to-one virtualization
    - Many physical machines/disks/networks may appear to look like one virtual machine/disk/network etc
  3. Many-to-many virtualization
    - Extend the above statements.


# Eg: Disk Virtualization



# Virtual Machines

- Logical/Emulated representations of full computing system environment
  - CPU + memory + I/O
  - Implemented by adding layers of software to the real machine to support the desired VM architecture.
- Uses:
  - Multiple OSes on one machine, including legacy OSes
  - Isolation
  - Enhanced security
  - Live migration of servers
  - Virtual environment for testing and development
  - Platform emulation
  - On-the-fly optimization
  - Realizing ISAs not found in physical machines

# Interfaces of a computer system



**Software**

**ISA**

**Hardware**

User ISA : 7



System ISA : 8

Syscalls : 3

ABI : 3,7



API : 2,7

# Machine Interfaces




# Types of VMs


- Process VMs
  - Virtualizes the ABI
  - Virtualization software = Runtime
    - Runs in non-privileged mode (user space)
    - Performs binary translation.
  - Terminates when guest process terminates
- System VMs
  - Virtualizes the ISA
  - Virtualization software = Hypervisor
    - Runs in privileged mode
    - Traps and emulates privileged instructions
  - Lasts as long as physical host is alive




# Uses of Virtualizing Software



(a)



(b)





(c)

- a) emulating one instruction set with another
- b) replicating a virtual machine so that multiple OSes can be supported simultaneously,
- c) composing virtual machine software to form a more complex, flexible system.


# Process Virtual Machines

- Process in a multiprogramming OS
  - Standard OS syscall interface + instruction set
  - Multiple processes, each with its own address space and virtual machine view.
- Emulators
  - Support one ISA on hardware designed for another ISA
  - Interpreter:
    - Fetches, decodes and emulates individual instructions. Slow.
  - Dynamic Binary Translator:
    - Blocks of source instructions converted to target instructions.
    - Translated blocks cached to exploit locality.




# Process Virtual Machines (cont'd)

- Same ISA Binary Optimizers
  - Optimize code on the fly
  - Same as emulators except source and target ISAs are the same.
- High-Level Language VMs
  - Virtual ISA (bytecode) designed for platform independence
  - Platform-dependent VM executes virtual ISA
  - E.g. Sun's JVM and Microsoft's CLI (part of .NET)
  - Both are stack-based VMs




# Hypervisor

- Also called Virtual Machine Monitor (VMM)
- A hypervisor is an operating system for operating systems
  - Provides a virtual execution environment for an entire OS and its applications
  - Controls access to hardware resources
  - When guest OS executes a privileged instruction, Hypervisor intercepts the instruction, checks for correctness and emulates the instruction.




# Type 1 Hypervisor (Classical System VMs)

- Hypervisor executes natively on the host ISA
- Hypervisor directly controls hardware and provides all device drivers
- Hypervisor emulates sensitive instructions executed by the Guest OS
- E.g. KVM and VMWare ESX Server



# Type-2 Hypervisors ( Hosted VMs)

- A host OS controls the hardware
- The Hypervisor runs partly in process space and partly in the host kernel
- Hypervisor relies on host OS to provide drivers
- E.g. VMWare Desktop Client



# Para-virtualized VMs


- Modify guest OS for better performance
  - Because “Trap and emulate” is expensive for many privileged instructions
- Para-virtualized VM sees a virtual hardware abstraction that is similar, but not identical to the real hardware.
- Guest OS is modified to replace some privileged instructions with “hypercalls” to the Hypervisor.
- Advantage: Results in lower performance overhead
- Disadvantage: Needs modification to the guest OS.
- E.g. Xen provides both para-virtual as well as full-virtualization
- Often traditional Hypervisors are partially para-virtualized
  - Device drivers in guest OS may be para-virtualized whereas CPU and Memory may be fully virtualized

# Examples

<b>Virtualizaton method</b>	<b>Type 1 hypervisor</b>	<b>Type 2 hypervisor</b>
Virtualization without HW support	ESX Server 1.0	VMware Workstation 1
Paravirtualization	Xen 1.0	
Virtualization with HW support	vSphere, Xen, Hyper-V	VMware Fusion, KVM, Parallels
Process virtualization		Wine

# Whole System VMs: Emulation


- Host and Guest ISA are different
- So emulation is required
- Hosted VM + emulation
- E.g. Virtual PC (Windows on MAC)



# Co-designed VMs

- The hypervisor is designed closely with (and possibly built into) a specific type of hardware ISA (or native ISA).
- Goal: Performance improvement of existing ISA (or guest ISA) during runtime.
- Hypervisor performs Emulation from Guest ISA to Native ISA.
- E.g. Transmeta Crusoe
  - Native ISA based on VLIW
  - Guest ISA = x86
  - Goal power savings

# Taxonomy




# Virtualizing individual resources in system VMs

# CPU Virtualization for VMs

- Each VM sees a set of “virtual CPUs”
- Hypervisors must emulate privileged instructions issued by guest OS.
- Modern ISAs provide special interfaces for Hypervisors to run VMs
  - Intel provides the VTx interface
  - AMD provides the AMD-v interface
- These special ISA interfaces allow the Hypervisors to efficiently emulate privileged instructions executed by the guest OS.
- When guest OS executes a privileged instruction
  - Hardware traps the instruction to the hypervisor
  - Hypervisor checks whether instruction must be emulated.
  - If so, Hypervisor reproduces the effect of the privileged operation.

# Execution of Privileged Instruction by Guest







Figure 8.4 Actions Taken by the VMM in Retiring One Virtual Machine and Activating the Next Virtual Machine

## Resource Control


- Issue: How to retain control of resources in the Hypervisor?
- Timer interval control performed by Hypervisor
- Also, guest OS is not allowed to read the timer value
  - Guest OS sees a virtual interval timer
- Hypervisor also gains control whenever guest OS executes privileged instructions

# Memory Virtualization for VMs

## Traditional virtual memory



## Virtual memory for VMs



- Guest OS in each VM sees a “guest”-physical address (GPA) space instead of the physical addresses
- Often hardware supports two-level page tables
  - EPT in Intel VT-x and NPT in AMD-v
- When hardware doesn’t, then Hypervisor needs to emulate two-level page tables using “shadow page tables”

# I/O Virtualization for VMs

- Hypervisor provides a virtual version of each physical device
- I/O activity directed at the virtual device is trapped by Hypervisor and converted to equivalent request for the physical device.
- Options:
  - Device emulation
    - Hypervisor traps and emulates each I/O instruction from Guest in Hypervisor.
    - Throughput is very slow.
    - Difficult to emulate the effect of combinations of I/O instructions.
  - Para-virtual devices
    - Special device drivers inserted in guest OS to talk to Hypervisor.
    - Most common.
  - Direct device access
    - Allow the VM to directly access physical device.
    - Fastest option but not scalable.
    - Requires IOMMU and VT-d support from hardware

Thank you!